

1 Project Documentation Template

1.1 Table of Contents

1. Introduction
 2. Project Structure
 3. Installation
 4. Usage
 5. Configuration
 6. Testing
 7. Deployment
 8. Contributing
 9. License
 10. Acknowledgments
 11. System Architecture
-

2 System's Perspective

2.1 Design and Architecture

2.2 Dependencies

Dependency List:

1. Microsoft.EntityFrameworkCore.Design - Version: 9.0.1
2. Microsoft.Extensions.Configuration - Version: 9.0.2
3. Microsoft.Extensions.Configuration.EnvironmentVariables - Version: 9.0.2
4. Microsoft.Extensions.Configuration.UserSecrets - Version: 9.0.2
5. Npgsql.EntityFrameworkCore.PostgreSQL - Version: 9.0.4
6. prometheus-net - Version: 8.2.1
7. Serilog.AspNetCore - Version: 9.0.0
8. Serilog.Sinks.Console - Version: 6.0.0
9. Microsoft.AspNetCore.Identity - Version: 2.3.1
10. Microsoft.EntityFrameworkCore.Sqlite - Version: 9.0.1
11. Microsoft.AspNetCore.Identity.EntityFrameworkCore - Version: 9.0.1
12. Microsoft.AspNetCore.Identity.UI - Version: 9.0.1
13. Microsoft.EntityFrameworkCore.Tools - Version: 9.0.0
14. Microsoft.VisualStudio.Web.CodeGeneration.Design - Version: 9.0.0
15. prometheus-net.AspNetCore - Version: 8.2.1
16. Serilog - Version: 4.2.0
17. Serilog.Formatting.Compact - Version: 3.0.0
18. Serilog.Sinks.Elasticsearch - Version: 10.0.0
19. Serilog.Sinks.Network - Version: 2.0.2.68
20. Serilog.Sinks.Async - Version: 1.5.0
21. coverlet.collector - Version: 6.0.4
22. Microsoft.AspNetCore.Mvc.Testing - Version: 9.0.2

23. Microsoft.NET.Test.Sdk - Version: 17.13.0
24. Microsoft.Playwright.NUnit - Version: 1.50.0
25. xunit - Version: 2.9.2
26. xunit.runner.visualstudio - Version: 3.0.0
27. Postgres - Version: 16.9
28. Kibana - Version: 8.12.1
29. logstash - Version: 8.12.1
30. elasticsearch - Version: 8.12.1
31. Nginx - Version: 1.27.0
32. Dotnet_SDK - Version: 9.0.0
33. org.Sonarcube - Version: 6.1.0

2.2.0.1 Logging For logging, our application uses Serilog as the API to collect log data. This data is then transferred into the Elastic Stack, which consists of Logstash, Elasticsearch, and Kibana—all used to process, query, and display the logging data. This setup is hidden behind Nginx, which acts as a reverse proxy and serves as an authentication layer between the user and Kibana (a data visualization and exploration tool). ##### Monitoring For monitoring, our application uses Prometheus as a real-time metrics storage server. On top of this, we use Grafana as a data visualization tool to display and analyze these metrics. ##### Application We have built our application using the .NET software framework, following the onion architecture originally invented by Jeffrey Palermo. We use the ASP.NET Core Identity package as an authentication system, allowing us to create and delete users. Initially, we used SQLite as our DBMS but later switched to Prometheus. In both cases, we utilized Entity Framework Core (EF Core) as our Object-Relational Mapper (ORM). For testing, we use NUnit as the primary testing framework, with Playwright layered on top for end-to-end testing. To handle API calls from the simulator, we use the ASP.NET Core MVC framework to create API controllers that process HTTP requests. As a software quality measure, we use SonarQube, specifically integrating their service via a GitHub workflow. SonarQube tracks security, reliability, maintainability, test coverage, and code duplications.

2.3 Interactions of Subsystems

2.3.0.1 Sequence Diagram for Simulator unfollow call

sequenceDiagram

```

    participant minitwit_Simulator
    participant Minitwit_Application
    participant DataBase

```

```

    minitwit_Simulator->>Minitwit_Application: Http Post (api/fllws/{username}) unfollow {us

```

```

    Note right of Minitwit_Application: UnFollowRequest

```

```

    Minitwit_Application-->>minitwit_Simulator: http statusCode 200

```

```

    Note left of minitwit_Simulator: Succesfull Response

```

```

    loop BatchInsert

```

```

Minitwit_Application->>Minitwit_Application: Insert into batch queue
end
Minitwit_Application->>DataBase: Get UserId <user1> <user2>
DataBase->>Minitwit_Application: UserId1 UserId2
Note left of Minitwit_Application: checks if both users exists
Minitwit_Application->>DataBase: Does User1 Follow User2
Note right of Minitwit_Application: FollowRequests User1 Follow User2
DataBase->>Minitwit_Application: Bollean
Note left of Minitwit_Application: If true
Note right of DataBase: Unfollow Sql Command
Minitwit_Application->>DataBase: Put user1 unfollow user2
DataBase->>Minitwit_Application: Status Response

```

2.3.0.2 Sequence Diagram for User unfollow call

```

sequenceDiagram
    participant User
    participant Minitwit_Application
    participant DataBase
    User->>Minitwit_Application: Http UnFollow(username)
    Note right of Minitwit_Application: UnFollowRequest
    Minitwit_Application->>DataBase: Get UserId <user1> <user2>
    DataBase->>Minitwit_Application: UserId1 UserId2
    Note left of Minitwit_Application: checks if both users exists
    Minitwit_Application->>DataBase: Does User1 Follow User2
    DataBase->>Minitwit_Application: Bollean
    Note left of Minitwit_Application: If true
    Note right of DataBase: Unfollow Sql Command
    Minitwit_Application->>DataBase: Put user1 unfollow user2
    DataBase->>Minitwit_Application: Status Response
    Minitwit_Application-->>User: http statuscode 200
    Note left of User: Succesfull Response

```

2.4 Current State of the System

3 Process' perspective

3.1 Deployment and Release

3.2 Monitoring

The application utilizes Prometheus and Grafana for monitoring. Prometheus scrapes port 5000, the minitwit application, and sends the data to the /metrics endpoint. Then Grafana retrieves the data from /metrics, and uses this as its data source. The relevant information could however not be found in the default configs . In the MetricsService.cs file, there are

custom metrics to our application, such as the “minitwit_follow_counter” and “app_request_duration_seconds”. The follow counter is implemented in the program by adding to the counter, every time a follow request is made. The duration is measured by starting a timer when a request comes in, and stopping it when the request has been processed.

3.3 Logging of application

The application uses the ELK logging stack. In the beginning, the logs contained information from the information level and up. This resulted in a flood of logs, and it was impossible to see anything relevant. It was then configured to only show warnings and above. Here there were practically no logs. From here logging statements were added to the code, to log when problems occurred. In the `ApiController.cs` there are custom creation of logs which are logged as warnings. These logs include system failures such as unsuccessful message post and failure to follow a user. This data is sent through Seriallog to Logstash. Another important metric is logging of request times. If a request took longer than 300 ms to process, it will log it. This has been central in discovering the ReadTimeout issue, that has occurred. To see all the logs for e.g. timeouts, the searchbar is used. Here the user can input “@m: slow”, to get them all.

3.4 Security assessment

3.5 Strategy for scaling and upgrade

Pay more money to digital ocean.

3.6 The use of AI

4 Reflection Perspective

4.1 Evolution and Refactoring

4.2 Operation

4.3 Maintenance

4.4 DevOps

5 Introduction

Provide a brief overview of the project, its purpose, and key features.

6 Project Structure

Explain the directory structure and the purpose of each folder/file.

```

/itu-minitwit
  .github/
    workflows
      build-and-test.yml      # GitHub Action workflows
      build-release.yml       # Automated build and test
      continous-deployment.yml # Creates release on push with a tag
      lint-and-format-check.yml # Deployment to dig
      scheduled-release.yml   # Automated linter and formatting checks
      sonarcube.yml           # Automated weekly release
                                # Automated Sonarcube checks
  logging/
    docker-compose.yml       # Logging configuration files
    nginx.conf               # Starts ELK stack and nginx containers
  logstash/
    remote_files/            # Reverse proxy with authentication
    report/                  # Logstash configuration
  src/
    minitwit.core/           # Files used remotely on the minitwit server for
    minitwit.infrastructure/ # Report files
    minitwit.web/            # Source code
      Program.cs              # Domain Layer - Domain models
  terraform/                 # Infrastructure Layer - Data access
    files/                   # Presentation Layer - Web app & API entry point
    modules/                 # Program entrypoin
      minitwit_logging/      # Terraform configurations for provisioning
      minitwit_server/       # Files used by terraform
      main.tf                # Terraform code for logging infrastucture
      terraform.tfvars       # Terraform code for minitwit infrastucture
      variables.tf           # Terraform module definitions
    tests/                   # Terraform variables
      minitwit.tests/        # Terraform variables declarations
      minitwit.tests.cs      # Test cases
      playwright.test.cs     # API tests
                                # UI tests

  docker-compose.yml         # For running the program locally
  Dockerfile                 # Application Dockerfile
  itu-minitwit.sln           # Project solution file

```

6.1 Installation

Step-by-step guide on how to set up the project locally.

Clone the repository

```
git clone https://github.com/Docker-Daze/itu-minitwit.git
```

Navigate to the project directory

```
cd itu-minitwit
```

```
# Install dependencies
dotnet restore
dotnet build

# Apply database migrations
dotnet ef database update
```

6.2 Usage

The deployed application is accessible on “http://164.90.240.84:5000/public”
How to deploy the application is mentioned in the deployment section.

The other ports running for monitoring and logging on the deployed app: -
Prometheus is on “164.90.240.84:9091” - Grafana is on “164.90.240.84:3000” -
elasticsearch is on “209.38.112.21:9200” - logstash is on “209.38.112.21:5044” -
kibana is on “209.38.112.21:5601” - nginx is on “209.38.112.21:8080”

It requires a login to access monitoring and logging.

For Grafana “164.90.240.84:3000” you can use the given login to access the dashboard. For elasticsearch “209.38.112.21:8080” use “admin” “admin” to login and access logs.

Run the application locally To run the application locally use this command inside the root folder itu-minutit:

```
# Command to run the project
docker compose up --build
```

When the minutitimage service is built and all containers are created, then the application is up and running on “localhost:5114/”. To access the application, navigate to “localhost:5114/” in the browser.

You can use the application's features like creating an account and post on the public timeline.

The ports for monitoring and logging are also accessible locally. - Prometheus is on “localhost:9091” - Grafana is on “localhost:3001” - elasticsearch is on “localhost:9200” - kibana is on “localhost:5601”

Note that no data and Grafana dashboard are set up locally.

6.3 Configuration

Details about configuration files and environment variables.

7 Process Perspective

This perspective should clarify how code or other artifacts come from idea into the running system and everything that happens on the way.

In particular, the following descriptions should be included:

7.1 Description of stages and tools included, including deployment and release of your systems.

7.2 Monitoring of application

The application utilizes Prometheus and Grafana for monitoring. Prometheus scrapes port 5000, the minitwit application, and sends the data to the /metrics endpoint. Then Grafana retrieves the data from /metrics, and uses this as its data source. The relevant information could however not be found in the default configs. In the `MetricsService.cs` file, there are custom metrics to our application, such as the “minitwit_follow_counter” and “app_request_duration_seconds”. The follow counter is implemented in the program by adding to the counter, every time a follow request is made. The duration is measured by starting a timer when a request comes in, and stopping it when the request has been processed.

7.3 Logging of application

The application uses the ELK logging stack. In the beginning, the logs contained information from the information level and up. This resulted in a flood of logs, and it was impossible to see anything relevant. It was then configured to only show warnings and above. Here there were practically no logs. From here logging statements were added to the code, to log when problems occurred. In the `ApiController.cs` there are custom creation of logs which are logged as warnings. These logs include system failures such as unsuccessful message post and failure to follow a user. This data is sent through Seriallog to Logstash. Another important metric is logging of request times. If a request took longer than 300 ms to process, it will log it. This has been central in discovering the ReadTimeout issue, that has occurred. To see all the logs for e.g. timeouts, the searchbar is used. Here the user can input “@m: slow”, to get them all.

7.4 Security assessment

7.5 Strategy for scaling and upgrade

Pay more money to digital ocean. - A complete description of stages and tools included in the CI/CD chains, including deployment and release of your systems. - How do you monitor your systems and what precisely do you monitor? - What do you log in your systems and how do you aggregate logs? - Brief results of the

security assessment and brief description of how did you harden the security of your system based on the analysis. - Applied strategy for scaling and upgrades.

7.6 CI/CD Chains

Our CI/CD Pipelines are built for the purpose of easier maintainance and security of our deployment to servers. ensuring that our infrastructure are up regardless of maintainance or failures for newly pushed commits.

```
/root
    workflows
        build-and-test.yml
        build-release.yml
        continous-deployment.yml
        lint-and-format-check.yml
        scheduled-release.yml
        sonarcube.yml
```

The following workflows are implemented to ensure a robust CI/CD pipeline:

1. **Build and Test Workflow**
This workflow automates the build process and runs all unit and integration tests to ensure code quality.
2. **Build Release Workflow**
Automatically creates a release when a new tag is pushed to the repository.
3. **Continuous Deployment Workflow**
Deploys the application to the production server upon successful completion of tests and builds.
4. **Lint and Format Check Workflow**
Ensures that the code adheres to the project's linting and formatting standards.
5. **Scheduled Release Workflow**
Automates weekly releases to ensure regular updates and maintenance.
6. **SonarQube Workflow**
Performs static code analysis using SonarQube to identify potential bugs and vulnerabilities.

Each workflow is defined in the `.github/workflows` directory and is triggered based on specific events such as pushes, pull requests, or scheduled intervals.

7.6.1 Deployment Chain

The deployment process follows a structured chain format to ensure reliability and minimize downtime. The steps are as follows:

1. Linting and Code Quality Checks

The code is first analyzed for adherence to linting and formatting standards. This ensures that the codebase remains clean and maintainable.

2. Integration Testing

Once the linting checks pass, the commit undergoes rigorous integration testing to validate that all components work together as expected.

3. Deployment with Rolling Updates

If the commit successfully passes all previous stages, the deployment process begins. Rolling updates are utilized to ensure a seamless transition. This approach guarantees that if the deployment encounters any issues, an unaffected backup server remains operational to handle the workload while the problem is resolved.

This deployment strategy ensures high availability and minimizes the risk of service disruption during updates.

```
# Run tests  
<insert test commands>
```

7.7 Deployment

Deploy the application

To deploy the application navigate to this folder:

```
# Folder  
/itu-minitwit/terraform
```

When inside folder run:

```
# Command initializes terraform files  
terraform init
```

Initialises terraform files if they are not already initialized Then run:

```
# Command show terraform changes  
terraform plan
```

This show what changes will be made when running terraform apply. finally run:

```
# Command apply terraform changes  
terraform apply  
# Confirm changes by saying yes  
yes
```

Wait for the application to deploy. When the application is deployed the website will be accessible on “<http://164.90.240.84:5000/public>”.

7.8 Contributing

Guidelines for contributing to the project.

7.9 License

This Project Itu_Minitwit is licensed and distributed under the MIT license

7.10 Acknowledgments

Credit individuals or resources that helped in the project.